

Measuring Context Switches in Serverless Environments

Elaine Yao
elainedv111@gmail.com

1 Introduction

Serverless computing is a new type of cloud application model. Instead of deploying virtual machines or containers, the developers only need to write functions without considering resource allocation, scheduling, etc. This simplifies the process of deploying code into production. Some commercial offerings include AWS lambda, Azure serverless and Google Cloud function. The cloud provider will allocate resources on demand and the pricing model of serverless platforms is different. It will bill the customers for the invocations times, functions execution time and the outbound network request. For example, in Google Cloud Function, when the amount of invocation is beyond 2 millions, for each extra million of invocation, it will charge \$0.4. For each gigabyte of outbound network, it will charge \$0.12. As for the execution time, the price is based on the amount memory and CPU resources provisioned for the function. In Table 1 For example, with 128 MB memory allocated and 0.2 GHz CPU frequency, the price for each second of execution time is \$0.00000231. And for \$1, the function deployed on cloud can run at most for a week. As the memory doubles, so does the price and CPU frequency. With more memory, the function will occupy more CPU resources. For exchange, the user has to pay more.

The execution time is measured from the time that the function receive a request, to the time it completes. It may complete through signaling, timeout or other failures. The problem is, user's functions can't be executed continuously due to limited resource. Providers want to invoke more functions at the same hardware to achieve the low resource cost. When resources like memory or CPU is limited, a running function may be interrupted to let another function being executed. Some CPU cycles are spent in OS scheduling, container management, page fault handling, etc. Thus, the function execution time(CPU wall-clock time) is larger than the actual execution time, which is presented in previous studies. [8] For users, who are charged with the function execution time, this phenomenon makes them charged more than they should be. In this project, we'll focus on the context switch in server,

which is one of the main extra time contributor [8].

Context switch [4] refers to the situation when the operating system interrupts the current execution and switch it off to another task in multithreading/processing. It happens when the number of processes/threads is more than the amount of CPU cores. In context switch, the OS has to save context registers for the currently-running process into the kernel stack, and then restore the context registers for the soon-to-be-executing process from kernel stack [2]. In this way, when the OS wants to resume the execution of previous process, it can just fetch the stored context registers.

There are various studies [3, 5–7] proposing benchmarks for measuring the context switch time in Linux systems. Most of them are creating pipes among different threads or processes and using system call to force the context switches. However, these benchmarks written in C can't be directly used in serverless environment as the latter only supports languages like Python, Go, PHP, Node.js, Java, etc. Also, there are other factors that may influence the number of the context switch and the context switch time in serverless environment, for example, the memory configurations. The characteristics of the context switch in the cloud are different from traditional Linux systems and thus the current benchmarks can't be used directly.

Have a sense of the context switch time in serverless environment can help user know about the extra cost they're paying for and may also motivate cloud providers to formulate more mature pricing model. In this work, we aim to answer the following research questions:

1. How to measure the context switch time accurately in a serverless environment?
2. What's the factors that may influence the number and time of the context switch in a serverless environment?

Memory	CPU frequency	Price/s	Time/\$
128 MB	0.2GHz	\$0.00000231	120h
256 MB	0.4GHz	\$0.00000463	60h
512 MB	0.8GHz	\$0.00000925	30h
1 G	1.4GHz	\$0.00001650	16.8h
2 G	2.4GHz	\$0.00002900	9.5h

Table 1: Context switch time by different benchmarks

2 Methodology

2.1 Challenges

Measuring the context switch time in serverless functions has several challenges:

C1 Characteristics of context switch in serverless environments

Context switch is triggered differently in serverless functions compared to in traditional operating systems as it's scheduled by the cloud provider. And the service provider has the tendency to trigger more functions in the same hardware to get more profits. Thus, context switch may happen not only due to performance factors like multi-threading/processing, but also provider's profit considerations.

C2 Benchmark accuracy

As none of the existing benchmarks can be used in serverless environment, it's challenging for us to reason about the accuracy of new benchmarks proposed. We also have to reason about the potential factors that might lead to variations in the measurement.

2.2 Reasoning context switches in serverless environment

In a cloud setting, the cloud provider will allocate multiple users to share the same physical device, in order to maximize the computing resource utilization. When there are multiple users co-locating at the same server, if the number of processes needed to run is larger than the cores inside the server, it's highly likely that there will be context switches.

Figure 1 shows one example scenario.

To tackle *C1*, we design the following experiments to analyze the factors influencing context switch in serverless environments. We measure the execution time elapse between two adjacent lines of code repeatedly under different configurations in serverless functions. Shahrade *et al.* [8] shows that the execution time in serverless functions is influenced by function invocation frequency, memory size and function execution time. Thus, these configurations will include the function invocation rate, memory allocated. Theoretically, the measured time elapse should be 0 as there is no execution

between these two lines. However, our initial experiments show that this is neither 0 nor a fixed value, indicating that interruption happens between these two lines and the time elapse changes with different configurations. Although the interruption is not limited to context switch but also page faults handling or container management, it can still provide insights in the influence of different configurations on context switch time in serverless environments.

2.3 Combine various benchmarks

To tackle *C2*, we first run our own designed benchmarks on local Linux system and compare the results with published benchmarks [1, 3, 6, 7], ensuring the correctness of the proposed ones. Then we deploy the new benchmarks in serverless functions and compare the results. If one of the benchmarks is consistently higher or lower than others, then we'll explore the reasons and this can help us improve our accuracy.

The previous benchmarks in context switching on Linux system mainly consider Our initial study to *C1* finds that context switch in serverless environments will also be influenced by invocation rate and memory allocation.

Currently, the new benchmarks are designed as the prior benchmarks with different configurations. The prior benchmarks we are choosing are:

1. Pingpong pipes [1, 7]: Two threads or processes are created. One thread/process writes the data to another process through a pipe, which induces a context switch. After the data is read, the read thread/process transfers the data to write thread/process and this induces another context switch. The total execution time equals the read and write operation time adds two context switch time. In contrast, a single process is created. It passes the data through a pipe to itself. Here the total execution time only consists of the read and write operation time. By subtracting these two time and dividing the result with 2, we get the context switch time.

2. Condition var [1]: A shared variable between two threads is used and is protected by mutex to prevent it from being modified by two threads at the same time. A signal is passed between two threads to tell which thread has the access to write the variable. And by passing the signals, the context will switch from one thread to the other. The time for the signal changes is the context switch time.

3. Lmbench [6]: A ring of pipes is created, and a token is passed from process to process with these pipes. Also, the time of passing the token in a single process with this ring of pipes is measured. Subtraction of these two time is the total context switch time of a number of context switches.

Next, we'll implement these functions with different invocation rate and memory allocation setting in the cloud, to measure the influence of configurations on the context switch time.




Figure 1: Context switches in serverless computing

3 Experiments

3.1 Experiment setup

We implement the benchmarks on Google Cloud, which is a mainstream cloud platform to provide serverless functions. For local measurements, we use an Ubuntu 20.04-64bit machine with Intel(R) Core(TM) i7-8750H CPU @ 2.20GHz processor and 12GB RAM. We plan to do experiments in more types of machines and operating systems to show the generality of the benchmark.

3.2 Measuring local context switches

3.2.1 Context switches time with different benchmarks

We write our own benchmark in Python, and we want to compare it with previous benchmarks to check its correctness. Therefore, we show the results for thread context switching for with pingpong method, condition var and process context switching with lmbench and pid sum, and perf (+pidstat/vmstat?) For lmbench, we choose xxx size. We run 100 times and get the median value and variance. And we draw them below in the figures. - bar plot, median and variance

1. Compare context switch between processes and threads.
2. Compare context switch between programming languages. (Pay attention to divide Python by 10)

3.2.2 Context switch time with different number of processes

I can't guarantee the size of the process

Benchmark	Pingpong	Condition Var	Lmbench	Pipe(python)
Time(/us)	1.96	2.29	1.78	33.6

Table 2: Context switch time by different benchmarks

Memory(/MB)	128	256	512
Time(/us)	81.11	33.96	36.17

Table 3: Context switch time in functions with different memories

3.3 Measuring context switches in the cloud

3.3.1 Non-parametric analysis among different memories

We first collect 200 data for each memory setting. And then we perform xxx method to check the normality. -> show the results for checking normality

Because it's not normal distribution, we then use non-parametric method to determine the number of experiment times.

We create functions with different memory allocated and measure the context switch time with the Ping pong pipes method. For each memory configuration we run 10 times and get the average of the calculated time. We also notice that the variation induced by Python language has already been considered and the value shown in Table 3 is processed. We notice that if the allocate memory is 128GB, the context switch time is remarkably higher than the other configurations. And we plan to look deeper into this in the future.

4 Future work

We plan to look at the following problems in the future.

1. Implement condition variable and lmbench in python and deploy the function in the cloud. We aim to measure the context switches in different ways in order to improve its accuracy.
2. Explore the factors that influence the number and the time of the context switch in serverless environment.

5 Related Work

A few studies [3, 5–7] have proposed benchmarks for measuring context switch in Linux system. Ousterhout *et al.* [7] proposes using two pipes to ping pong data between two processes. McVoy *et al.* [6] reduces the cost of system calls for

context switch by designing a ring of pipes, and they also consider the influence of data size and cache size on the context switch time [5]. Suo *et al.* [9] measures the context switch on the edge in AI workloads. None of them considers the context switch in a serverless environment.

References

- [1] Measuring context switching and memory overheads for linux threads. <https://eli.thegreenplace.net/2018/measuring-context-switching-and-memory-overheads-for-linux-threads/>. Accessed: 2022-2-25.
- [2] Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau. *Operating Systems: Three Easy Pieces*. Arpaci-Dusseau Books, 1.00 edition, August 2018.
- [3] Francis M. David, Jeffrey C. Carlyle, and Roy H. Campbell. Context switch overheads for linux on arm platforms. In *Proceedings of the 2007 Workshop on Experimental Computer Science, ExpCS '07*, page 3–es, New York, NY, USA, 2007. Association for Computing Machinery.
- [4] David Kalinsky. Context switch. *Embedded Systems Programming*, 14(1):94–105, 2001.
- [5] Chuanpeng Li, Chen Ding, and Kai Shen. Quantifying the cost of context switch. In *Proceedings of the 2007 Workshop on Experimental Computer Science, ExpCS '07*, page 2–es, New York, NY, USA, 2007. Association for Computing Machinery.
- [6] Larry W. McVoy and Carl Staelin. Imbench: Portable tools for performance analysis. In *USENIX Annual Technical Conference*, 1996.
- [7] John K. Ousterhout. Why aren't operating systems getting faster as fast as hardware? In *USENIX Summer*, 1990.
- [8] Mohammad Shahrad, Jonathan Balkind, and David Wentzlaff. Architectural implications of function-as-a-service computing. *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture*, 2019.
- [9] Kun Suo, Yong Shi, Chih-Cheng Hung, and Patrick Bobbie. *Quantifying Context Switch Overhead of Artificial Intelligence Workloads on the Cloud and Edges*, page 1182–1189. Association for Computing Machinery, New York, NY, USA, 2021.